


# From Monolith to Micro-services with Kubernetes

16 Mar 2019, FOSS Asia, Singapore



Michael Bright,  @mjbright

Slides & source code at <https://mjbright.github.io/Talks>

# TEST PAGE



**TL-Nature**

What a beautiful sunrise

**BL-Nature**

What a beautiful sunrise



**CM-Nature**

What a beautiful sunrise

**TR-Nature**

What a beautiful sunrise

**BR-Nature**

What a beautiful sunrise

Michael Bright,  @mjbright

Freelance Trainer: Kubernetes, Serverless, Docker,  
CloudNative

Past researcher, dev, team lead, dev advocate

British, living in France for 27-years

Docker Community Lead, Python User Group



[linkedin.com/in/mjbright](https://www.linkedin.com/in/mjbright)



[github.com/mjbright](https://github.com/mjbright)


# Outline

- Monoliths to Micro-services
- Orchestration: Kubernetes
- Deployment Strategies
- Architecture Design patterns
- Summary

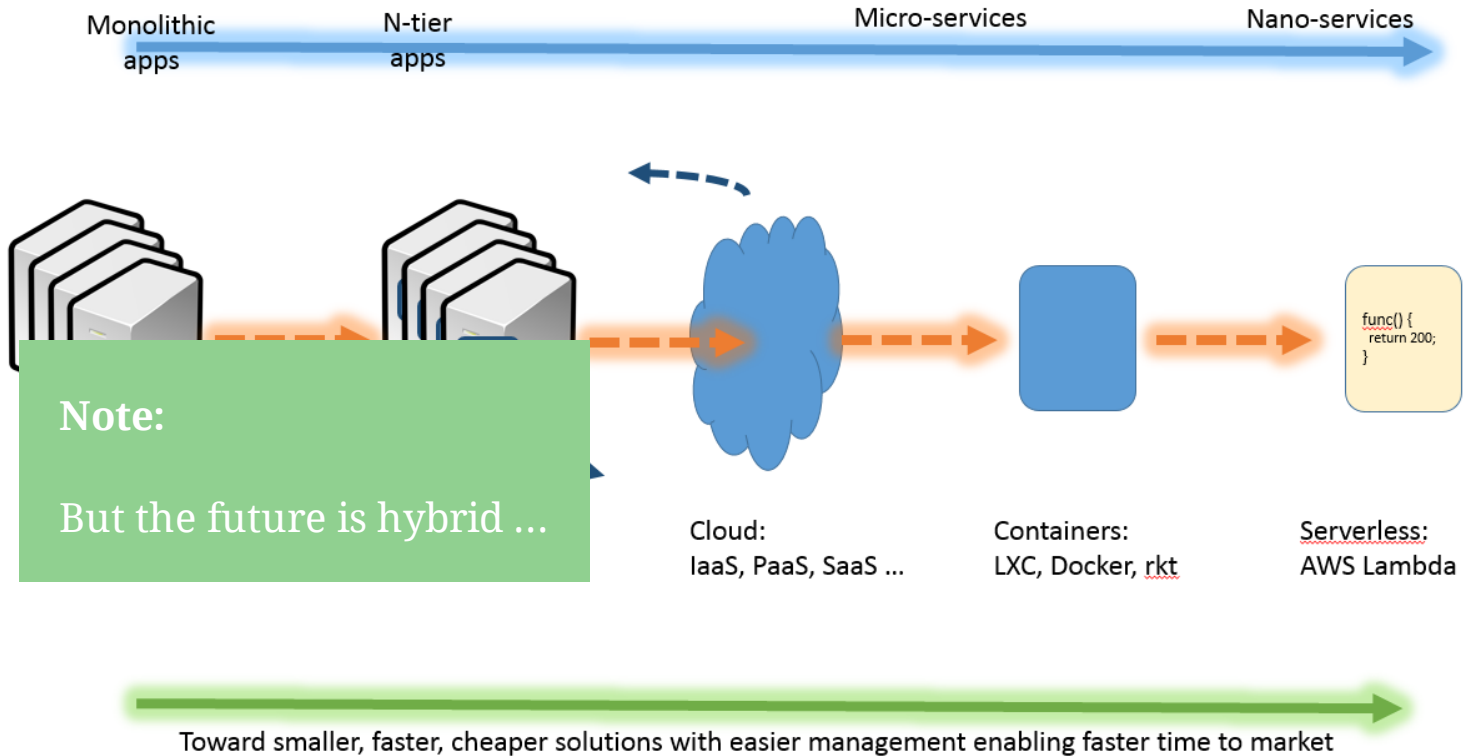
# Outline

- Monoliths to Micro-services
- Orchestration: Kubernetes
- Deployment Strategies
- Architecture Design patterns
- Summary

# First ... a bit of history




# First ... a bit of history



# Monoliths to Micro-services

Monoliths are **deployed, scaled, upgraded, reimplemented** as complete units



## Problem: Feel the pain


- components are tightly coupled and only scale as a unit.
- are developed in Waterfall
- difficult to patch
- library dependencies between components
- difficult to reuse



# Monoliths to Micro-services

## Proposition: split up components

Individual  $\mu$ -service components can be **deployed, scaled, upgraded, reimplemented ...**



In micro-service architecture components are only lightly coupled

- interconnected by network
- can be scaled independently
- can be deployed/upgraded independently

# Advantages of Micro-services

Separation of Concerns - "do one thing well"

# Advantages of Micro-services

Separation of Concerns - "do one thing well"

Smaller Projects/teams

# Advantages of Micro-services

Separation of Concerns - "do one thing well"

Smaller Projects/teams

Ease Scaling, Deployment, Testing, Evolution

# Advantages of Micro-services

Separation of Concerns - "do one thing well"

Smaller Projects/teams

Ease Scaling, Deployment, Testing, Evolution

Loosely coupled components

# Advantages of Micro-services

Separation of Concerns - "do one thing well"

Smaller Projects/teams

Ease Scaling, Deployment, Testing, Evolution

Loosely coupled components

Allow for composition of new services

Can be re-implemented

# Advantages of Micro-services

Separation of Concerns - "do one thing well"

Smaller Projects/teams

Ease Scaling, Deployment, Testing, Evolution

Loosely coupled components

Allow for composition of new services

Can be re-implemented

So are they a panacea?

# Disadvantages

## Greater complexity

- Requires orchestration, component version management
- Greater organizational complexity
- Monitoring, debugging, end-2-end test are more difficult



# Disadvantages

## Greater complexity

- Requires orchestration, component version management
- Greater organizational complexity
- Monitoring, debugging, end-2-end test are more difficult

## More network communication

- Network error handling, Performance, Circuit-breakers

# Disadvantages

## Greater complexity

- Requires orchestration, component version management
- Greater organizational complexity
- Monitoring, debugging, end-2-end test are more difficult

## More network communication

- Network error handling, Performance, Circuit-breakers

## Still requires best practices

- Behaviour and Test-Driven Development, CI/CD
- Documentation of interfaces/APIs, Stable

# Outline

- Monoliths to Micro-services
- Orchestration: Kubernetes
- Deployment Strategies
- Architecture Design patterns
- Summary

# Orchestration: Kubernetes

**Problem: Feel the pain** Impossible to manage 10000 containers running across a data center of 1000 nodes.

- on which nodes should you schedule
- which containers are malfunctioning
- which are started and ready to go ==>



# We need Orchestration

REPLACE: As you move from deploying containers on a single machine to deploying them across a number of machines, you'll need an orchestration tool to manage (and automate) the arrangement, coordination, and availability of the containers across the entire system.

REPLACE: Issues:

- Cross-server container communication
- Horizontal scaling
- Service discovery
- Security/TLS
- Zero-downtime deploys
- Rollbacks






# REPLACE: Orchestration Feature Wish-list

Feature Info Health checks Verify when a task is ready to accept traffic  
Dynamic port-mapping Ports are assigned dynamically when a new container is spun up  
Zero-downtime deployments Deployments do not disrupt end users  
Service discovery Automatic detection of new containers and services  
Auto scaling Automatically scale resources up or down based on the load  
Provisioning New containers should select hosts based on resources and configuration

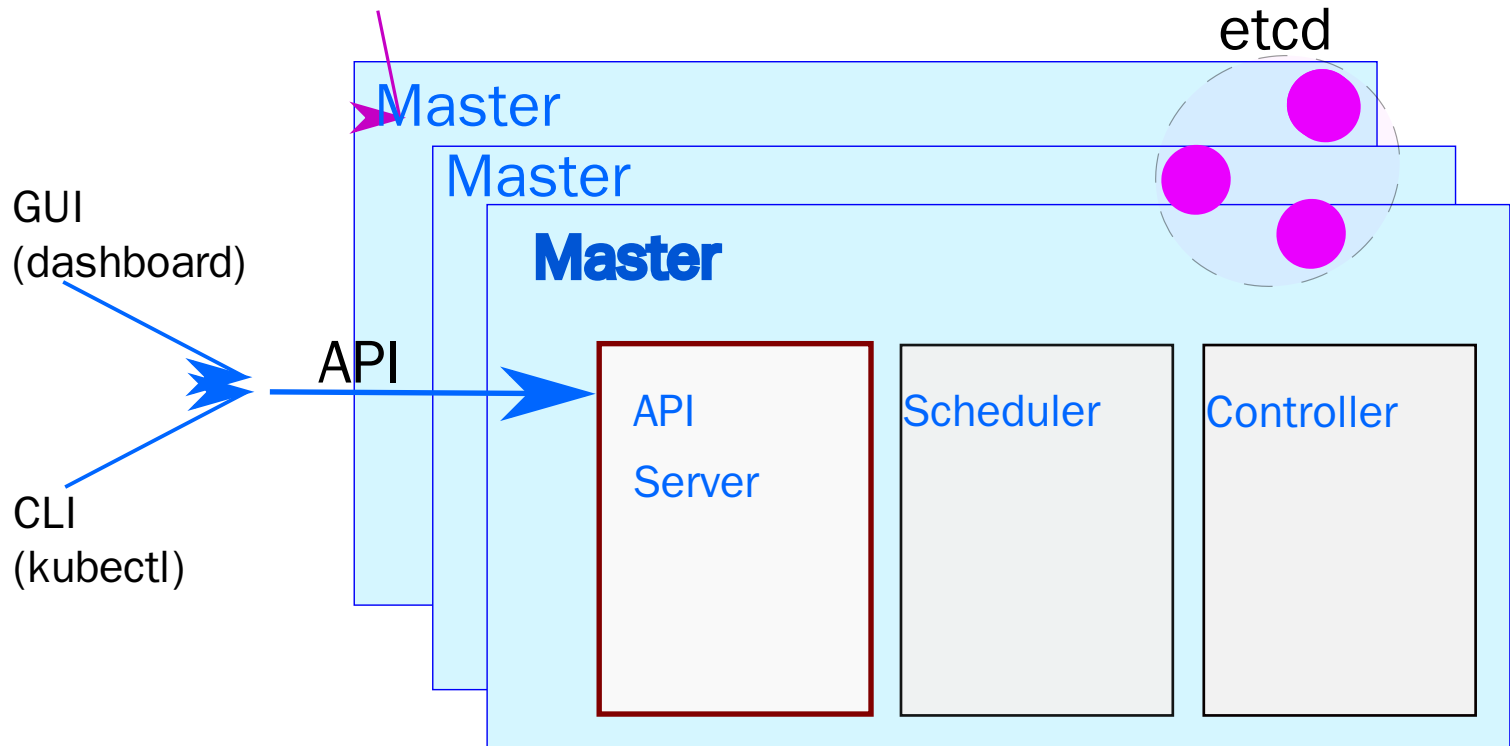
REPLACE: What else?

Load balancing, logging, monitoring, authentication and authorization, security... predictability, scalability, and high availability...

# Kubernetes - Architecture




# Kubernetes - Master Nodes





# Kubernetes - Worker Nodes



# Kubernetes - Pods

Containers share some namespaces:  
- PID, IPC, network , time sharing

Main container



Sidecar

Sidecar


same ip, e.g. 192.168.1.20

A pod houses one or more containers

# Kubernetes Demo



# Kubernetes - Deploying Redis



# Kubernetes - Deploying Redis

```
# kubectl run redis --image=redis:latest --port=6379
```

```
$ kubectl apply -f redis-deployment.yaml  
deployment.extensions "redis" created
```


```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
redis-68595c4d95-rr4pr	0/1	ContainerCreating	0	1s

# Kubernetes - Deploying Redis (yaml)

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  labels:
    run: redis
  name: redis
spec:
  replicas: 1
  selector:
    matchLabels:
      run: redis
  template:
    metadata:
      labels:
        run: redis
    spec:
      containers:
        - image: redis:latest
          name: redis
          ports:
            - containerPort: 6379
```

# Kubernetes - Deploying Flask



# Kubernetes - Deploying Flask

```
# kubectl run flask-app --image=$IMAGE --port=5000
```

```
$ kubectl apply -f flask-deployment.yaml  
deployment.extensions "flask-app" created
```

```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
flask-app-8577b44db-96cht	0/1	Pending	0	1s
redis-68595c4d95-rr4pr	0/1	ContainerCreating	0	1s



# Kubernetes - Deploying Flask (yaml)


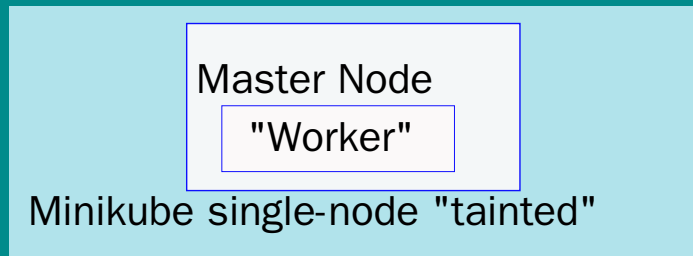
```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  labels:
    run: flask-app
  name: flask-app
spec:
  replicas: 1
  selector:
    matchLabels:
      run: flask-app
  template:
    metadata:
      labels:
        run: flask-app
    spec:
      containers:
        - image: mjbright/flask-web:v1
          name: flask-app
          ports:
            - containerPort: 5000
```

# Operations - Scaling

```
# kubectl scale deploy flask-app --replicas=4  
$ kubectl edit -f flask-deploy.yaml
```

```
...  
spec:  
  replicas: 4
```

# Demo



# Outline

- Monoliths to Micro-services
- Orchestration: Kubernetes
- Deployment Strategies
- Architecture Design patterns
- Summary

????  
....

**Problem: Feel the pain**

... so ... **Proposition: split up components**

# Micro-service Deployment Strategies

Rolling Upgrade

Health Checks

Strangler Pattern

# Operations - Rolling Upgrades

Several strategies exist

**recreate** - terminate old version before releasing new one

# Operations - Rolling Upgrades

Several strategies exist

**recreate** - terminate old version before releasing new one

**ramped** - gradually release a new version on a rolling update fashion



# Operations - Rolling Upgrades

Several strategies exist

**recreate** - terminate old version before releasing new one

**ramped** - gradually release a new version on a rolling update fashion

**blue/green** - release new version alongside old version then switch

# Operations - Rolling Upgrades

Several strategies exist

**recreate** - terminate old version before releasing new one

**ramped** - gradually release a new version on a rolling update fashion

**blue/green** - release new version alongside old version then switch

**canary** - release new version to subset of users, proceed to full rollout

# Operations - Rolling Upgrades

Several strategies exist

**recreate** - terminate old version before releasing new one

**ramped** - gradually release a new version on a rolling update fashion

**blue/green** - release new version alongside old version then switch

**canary** - release new version to subset of users, proceed to full rollout

**a/b testing** - release new version to subset of users in a precise way (HTTP headers, cookie, weight, etc.).



# Operations - Rolling Upgrade

## Ramped

```
# kubectl set image deploy flask-app flask-app=mjbright/flask-web:v2  
  
$ kubectl edit -f flask-deploy.yaml  
$ kubectl rollout status deployment/flask-app
```

```
...  
  spec:  
    containers:  
      - image: mjbright/flask-web:v2
```

# Demo



????  
....

**Problem: Feel the pain**

... so ... **Proposition: split up components**

# Operations - Healthchecks

Cover healthchecks  
and Readiness probes

????  
....

**Problem: Feel the pain**

... so ... **Proposition: split up components**



# Operations - Strangler Pattern

The Strangler is a pattern used in the initial migration from a Monolithic architecture to a Micro-services architecture

# Micro-service - Architecture Design Patterns

We are not concerned with:

Standard Component Patterns

Micro-services themselves (!) - Fine-grained SOA

Sidecar

# Micro-service - Architecture Design Patterns

We are concerned with:

Exposing Services

Ingress

API Gateway

Service Mesh


Hybrid Apps

????  
....

**Problem: Feel the pain**

... so ... **Proposition: split up components**

# Kubernetes - Exposing Services



# Design Pattern - Services

Services can be exposed via


NodePort

HostPort


ClusterIP

LoadBalancer

# Exposing Services (LoadBalancer)



# Exposing Services (NodePort)






????  
....

**Problem: Feel the pain**

... so ... **Proposition: split up components**

# Exposing Services (IngressController)



# Exposing Redis Service (LoadBalancer)

```
# kubectl expose deployment redis --type=LoadBalancer
```

```
$ kubectl apply -f redis-service.yaml  
service "redis" created
```

```
$ kubectl get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	5h
redis	LoadBalancer	10.101.158.201	<pending>	6379:31218/TCP	1s

# Exposing Redis Service (LoadBalancer)

```
apiVersion: v1
kind: Service
metadata:
  labels:
    run: redis
  name: redis
spec:
  ports:
    - port: 6379
      protocol: TCP
      targetPort: 6379
  selector:
    run: redis
  type: LoadBalancer
```

# Exposing Flask Service (LoadBalancer)

```
# kubectl expose deployment flask-app --type=LoadBalancer
```

```
$ kubectl apply -f flask-service.yaml  
service "flask-app" created
```


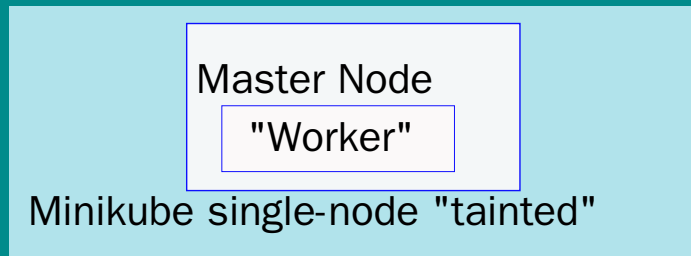
```
$ kubectl get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
flask-app	LoadBalancer	10.103.154.19	<pending>	5000:32201/TCP	1s
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	5h
redis	LoadBalancer	10.101.158.201	<pending>	6379:31218/TCP	2s

# Exposing Flask Service (LoadBalancer)

```
apiVersion: v1
kind: Service
metadata:
  labels:
    run: flask-app
  name: flask-app
spec:
  ports:
    - port: 5000
      protocol: TCP
      targetPort: 5000
  selector:
    run: flask-app
  type: LoadBalancer
```

# Demo



# Design Pattern - Ingress

## Ingress

Relation to API GW, LoadBalancer

## Ingress Rules

## Ingress Controller

## Ingress Gateway ?



# Exposing Services (Ingress)

```
$ minikube addons enable ingress  
ingress was successfully enabled  
  
$ kubectl apply -f misc/ingress-definition.yaml  
ingress.extensions "ingress-definitions" created  
  
$ sudo vi /etc/hosts  
...  
192.168.99.100 minikube.test flaskapp.test
```

# Exposing Services (Ingress)

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: ingress-definitions
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
spec:
  backend:
    serviceName: default-http-backend
    servicePort: 80
  rules:
  - host: minikube.test
    http:
      paths:
      - path: /
        backend:
          serviceName: k8sdemo
          servicePort: 8080
  - host: flaskapp.test
    http:
      paths:
      - path: /flask
        backend:
          serviceName: flask-app
          servicePort: 5000
```

# Exposing Services (Ingress)

```
$ minikube service list
```

NAMESPACE	NAME	URL
default	flask-app	http://192.168.99.100:32201
default	k8sdemo	http://192.168.99.100:31280
default	redis	http://192.168.99.100:31218
kube-system	kubernetes-dashboard	http://192.168.99.100:30000

```
$ curl http://192.168.99.100:31280
```

```
$ curl http://minikube.test/k8sdemo
```

# Exposing Services (Ingress)

```
$ minikube service list
```

NAMESPACE	NAME	URL
default	flask-app	http://192.168.99.100:32201
default	k8sdemo	http://192.168.99.100:31280
default	redis	http://192.168.99.100:31218
kube-system	kubernetes-dashboard	http://192.168.99.100:30000



```
$ curl http://192.168.99.100:32201
```

```
[flask-app-8577b44db-kbwpm] Redis counter value=214
```

```
$ curl http://flaskapp.test/flask
```

```
[flask-app-8577b44db-kbwpm] Redis counter value=215
```

# Demo




????  
....

**Problem: Feel the pain**

... so ... **Proposition: split up components**


# Design Pattern - API Gateway



External endpoint exposes APIs

- Offloads common Ingress functions => reduces  $\mu$ -service complexity
  - rate limiting, security, authorisation, DDOS protection
  - Protocol version translation, e.g. REST to SOAP, \*-RPC ...
  - TLS decryption/encryption
- Hides internal infrastructure detail => controls access
  - service routing, load-balancing
  - Allows to refactor/scale/mock internal implementation

# Design Pattern - API Gateway



External endpoint exposes APIs

- Offloads common Ingress functions => reduces  $\mu$ -service complexity
  - rate limiting, security, authorisation, DDOS protection
  - Protocol version translation, e.g. REST to SOAP, \*-RPC ...
  - TLS decryption/encryption
- Hides internal infrastructure detail => controls access
  - service routing, load-balancing
  - Allows to refactor/scale/mock internal implementation

Needs to scale, be H.A.



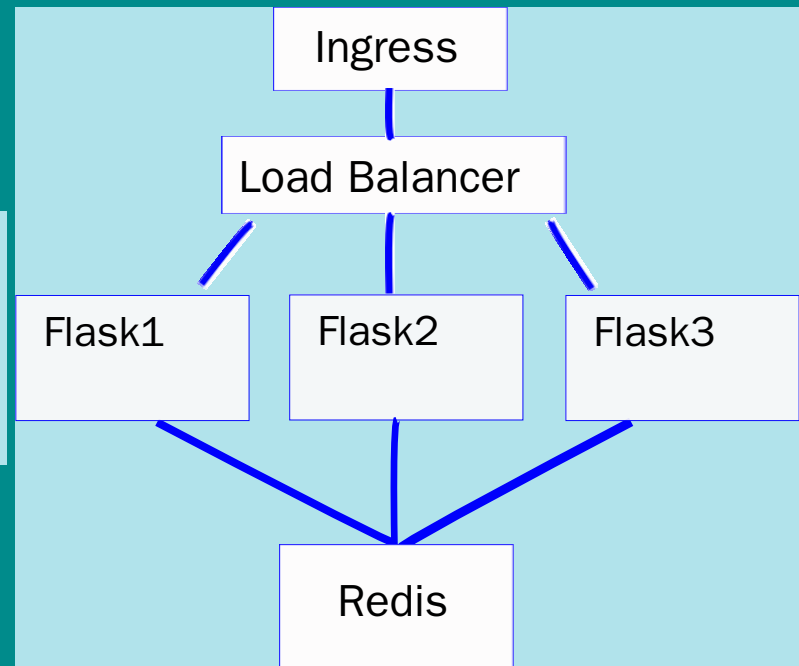

# Design Pattern - API Gateway

There are many API Gateways including

ha-proxy, NGInx, Traefik, AWS ELB ?!

Newer generation: Envoy-based such as Ambassador, Gloo

# Demo



????  
....

**Problem: Feel the pain**

... so ... **Proposition: split up components**

# Design Pattern - Service Mesh

Abstraction above TCP/IP, secure reliable inter-service connectivity.

Platforms such as Linkerd (v2) and Istio (v1) provide offload for  $\mu$ -services

# Design Pattern - Service Mesh

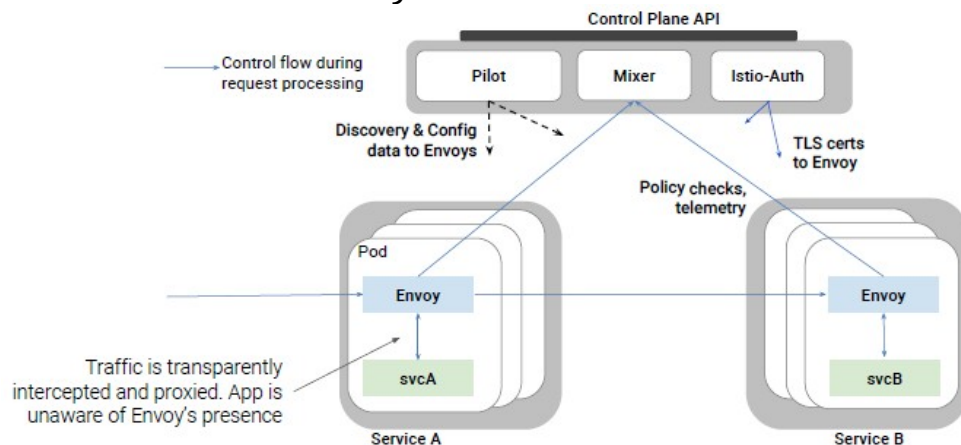
Abstraction above TCP/IP, secure reliable inter-service connectivity.

Platforms such as Linkerd (v2) and Istio (v1) provide offload for  $\mu$ -services  
Offloads functionality from services in a distributed way.

# Design Pattern - Service Mesh

Abstraction above TCP/IP, secure reliable inter-service connectivity.

Platforms such as Linkerd (v2) and Istio (v1) provide offload for  $\mu$ -services  
Offloads functionality from services in a distributed way.



# Design Pattern - Service Mesh - Linkerd

Abstraction above TCP/IP, secure reliable inter-service connectivity.

Platforms such as Linkerd (v2) and Istio (v1) provide offload for  $\mu$ -services

# Design Pattern - Service Mesh - Linkerd

Abstraction above TCP/IP, secure reliable inter-service connectivity.

Platforms such as Linkerd (v2) and Istio (v1) provide offload for  $\mu$ -services

Offloads functionality from services in a distributed way.

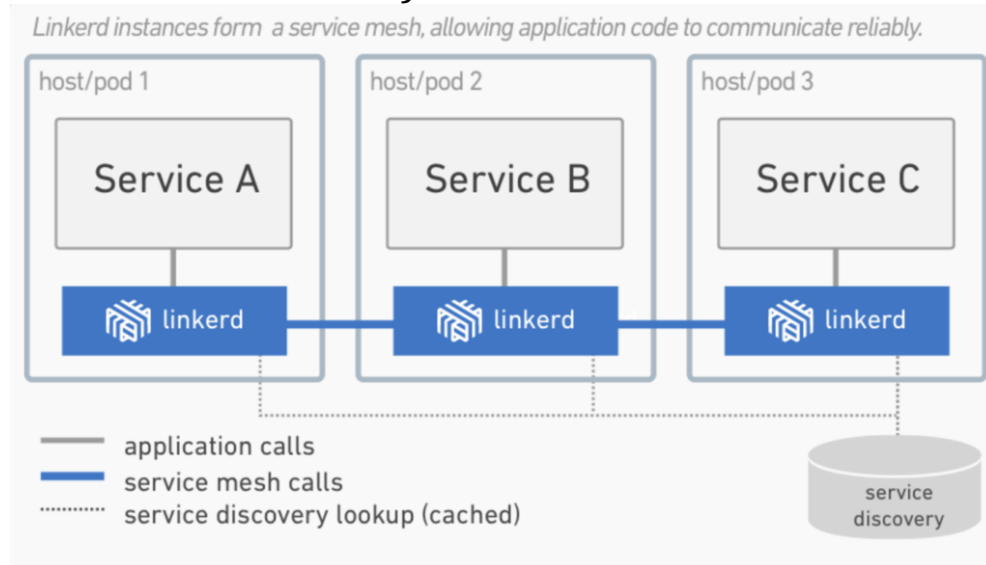


# Design Pattern - Service Mesh - Linkerd

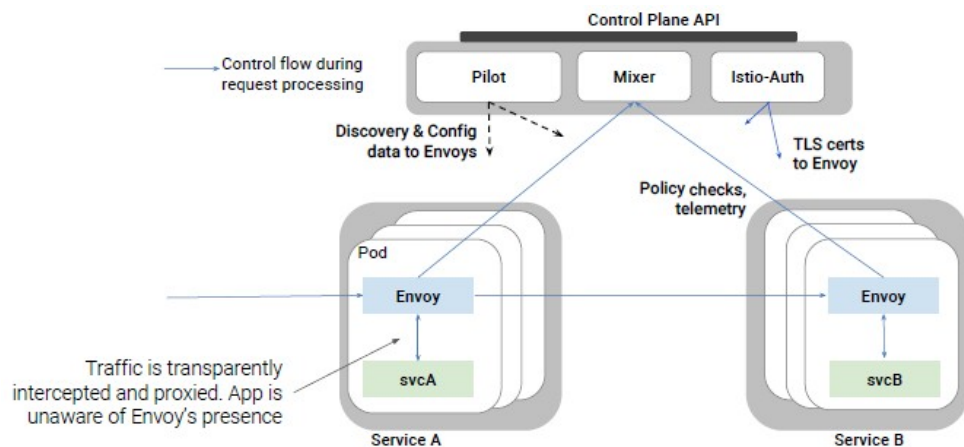
Abstraction above TCP/IP, secure reliable inter-service connectivity.

Platforms such as Linkerd (v2) and Istio (v1) provide offload for  $\mu$ -services



Offloads functionality from services in a distributed way.



# Design Pattern - Service Mesh - Istio



# Demo




????  
....

**Problem: Feel the pain**

... so ... **Proposition: split up components**

# Design Pattern - Hybrid Apps

Gloo allows to route between legacy apps, micro-services and serverless incrementally adding new functionality.

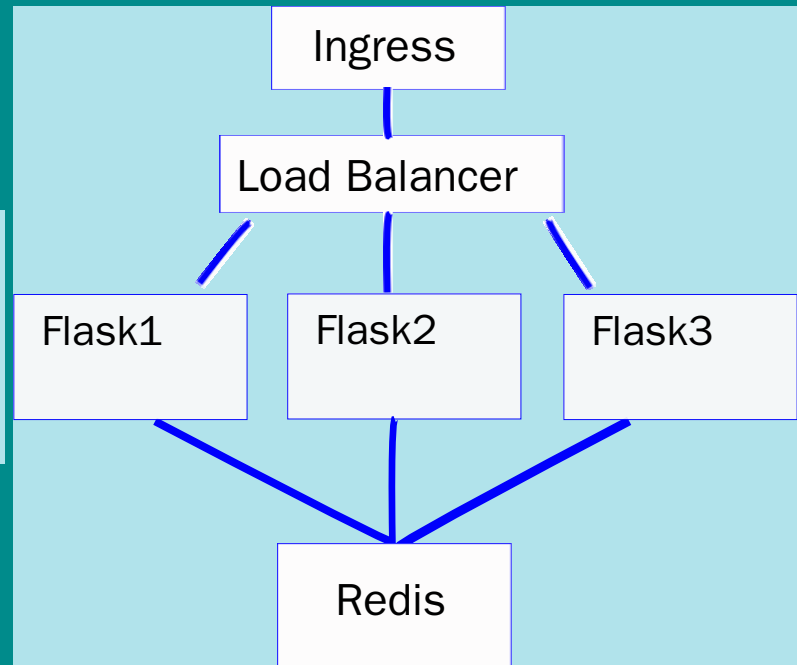



<https://medium.com/solo-io/building-hybrid-apps-with-gloo-1eb96579b070>

# Design Pattern - Hybrid Apps

Gloo understands the infrastructure on which it is running and the APIs being used.

# Demo



# Outline

- Monoliths to Micro-services
- Orchestration: Kubernetes
- Deployment Strategies
- Architecture Design patterns
- Summary



????  
....

**Problem: Feel the pain**

... so ... **Proposition: split up components**

# Tools

WHAT / WHERE ????

- Tools
  - Helm (use to install tools)
  - Prometheus
  - Squash
  - Gloo
  - Istio / Service Meshes / Envoy



# Summary

Micro-services offer new deployment possibilities

- with ease of deployment, scaling, upgrading
- facilitate "Best in Class" technology choices/replacements

# Summary

Micro-services offer new deployment possibilities

- with ease of deployment, scaling, upgrading
- facilitate "Best in Class" technology choices/replacements

*BUT* moving to  $\mu$ -services requires

- organizational changes and best practices !
- incremental rollout - small steps / Strangler
- hybrid approaches - old/new, cloud/on-premise, VM/container/ $\mu$ -service
- offload via API Gateway and/or Service Mesh

# Thank you !

From Monologue to Discussions ... ?

## Questions ?

Michael Bright,  @mjbright

Cloud Native Training (Docker, Kubernetes, Serverless)

Slides & source code at <https://mjbright.github.io/Talks>



[linkedin.com/in/mjbright](https://www.linkedin.com/in/mjbright)  [github.com/mjbright](https://github.com/mjbright)

# Summary

## Getting started with Kubernetes

Start by learning Docker principles

Experiment by Dockerizing some applications

Learn about Container Orchestration

Hands-on with Kubernetes online or  
Minikube(\*)

Kubernetes Visualization with KubeView

<https://github.com/mjbright/kubeview>

# Resources



**minikube**

- Download <https://github.com/kubernetes/minikube/releases>
- Documentation <https://kubernetes.io/docs/getting-started-guides/minikube/>
- Hello Minikube <https://kubernetes.io/docs/tutorials/stateless-application/hello-minikube/>

# Resources - Articles

Martin Fowler	<a href="https://martinfowler.com/articles/microservices.html">https://martinfowler.com/articles/microservices.html</a>
MuleSoft, "The top 6 Microservices Patterns"	<a href="https://www.mulesoft.com/lp/whitepaper/api/top-microservices-patterns">https://www.mulesoft.com/lp/whitepaper/api/top-microservices-patterns</a>
FullStack Python	<a href="https://www.fullstackpython.com/microservices.html">https://www.fullstackpython.com/microservices.html</a>
Idit Levine	<a href="https://medium.com/solo-io/building-hybrid-apps-with-gloo-1eb96579b070">https://medium.com/solo-io/building-hybrid-apps-with-gloo-1eb96579b070</a>
SSola	<a href="https://medium.com/@ssola/building-microservices-with-python-part-i-5240a8dcc2fb">https://medium.com/@ssola/building-microservices-with-python-part-i-5240a8dcc2fb</a>
Deployment	<a href="http://container-solutions.com/kubernetes-deployment-strategies/">http://container-solutions.com/kubernetes-deployment-strategies/</a>

TO ADD: 12-factor apps Gloo - Christian Posta



# Resources - Books

## Publisher

O'Reilly



PacktPub



## Title, Author

"Building Microservices", Sam Newman,  
July 2015

"Python Microservices Development",  
**Tarek Ziade**, July 2017

kNative - O'Reilly

Istio - Manning

Istio - O'Reilly

Testdriven.io